

Assignment #1

Implementing a network service that can handle large scale network requests

** Co-worked with Jaehyun Nam*

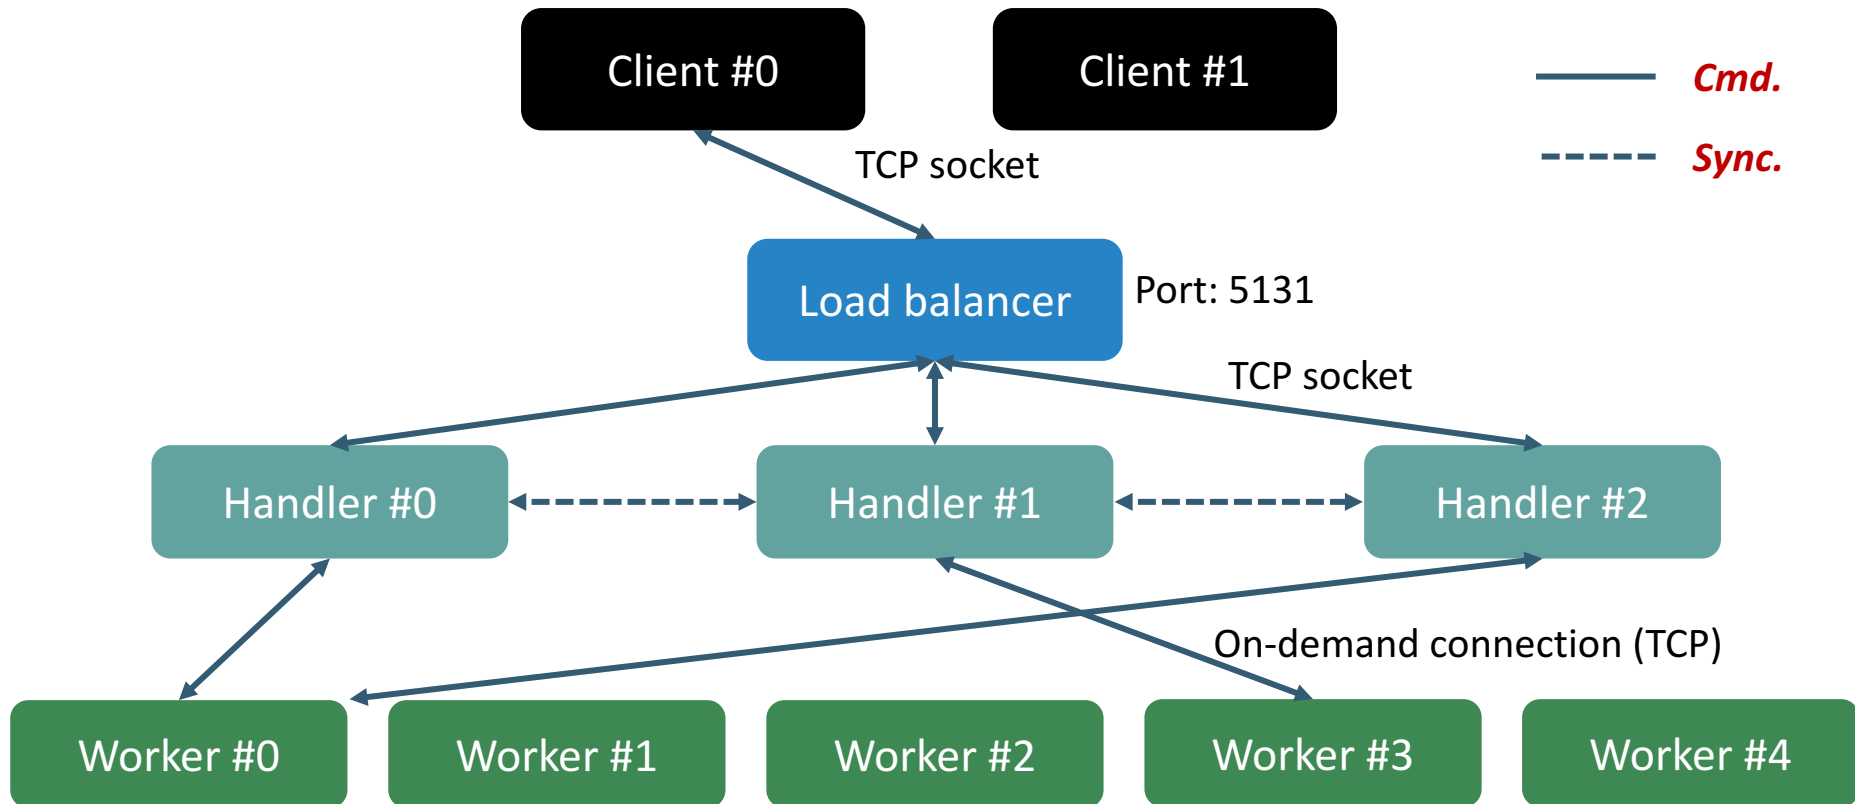
High-level functions

- *Ultimate Goal: Build a distributed key-value storage system*
- Store/retrieve data as a *key-value pair*
- Handle *bulk requests* with I/O multiplexing and load balancing
- Write *logs* and provide *CLIs* to interact with users



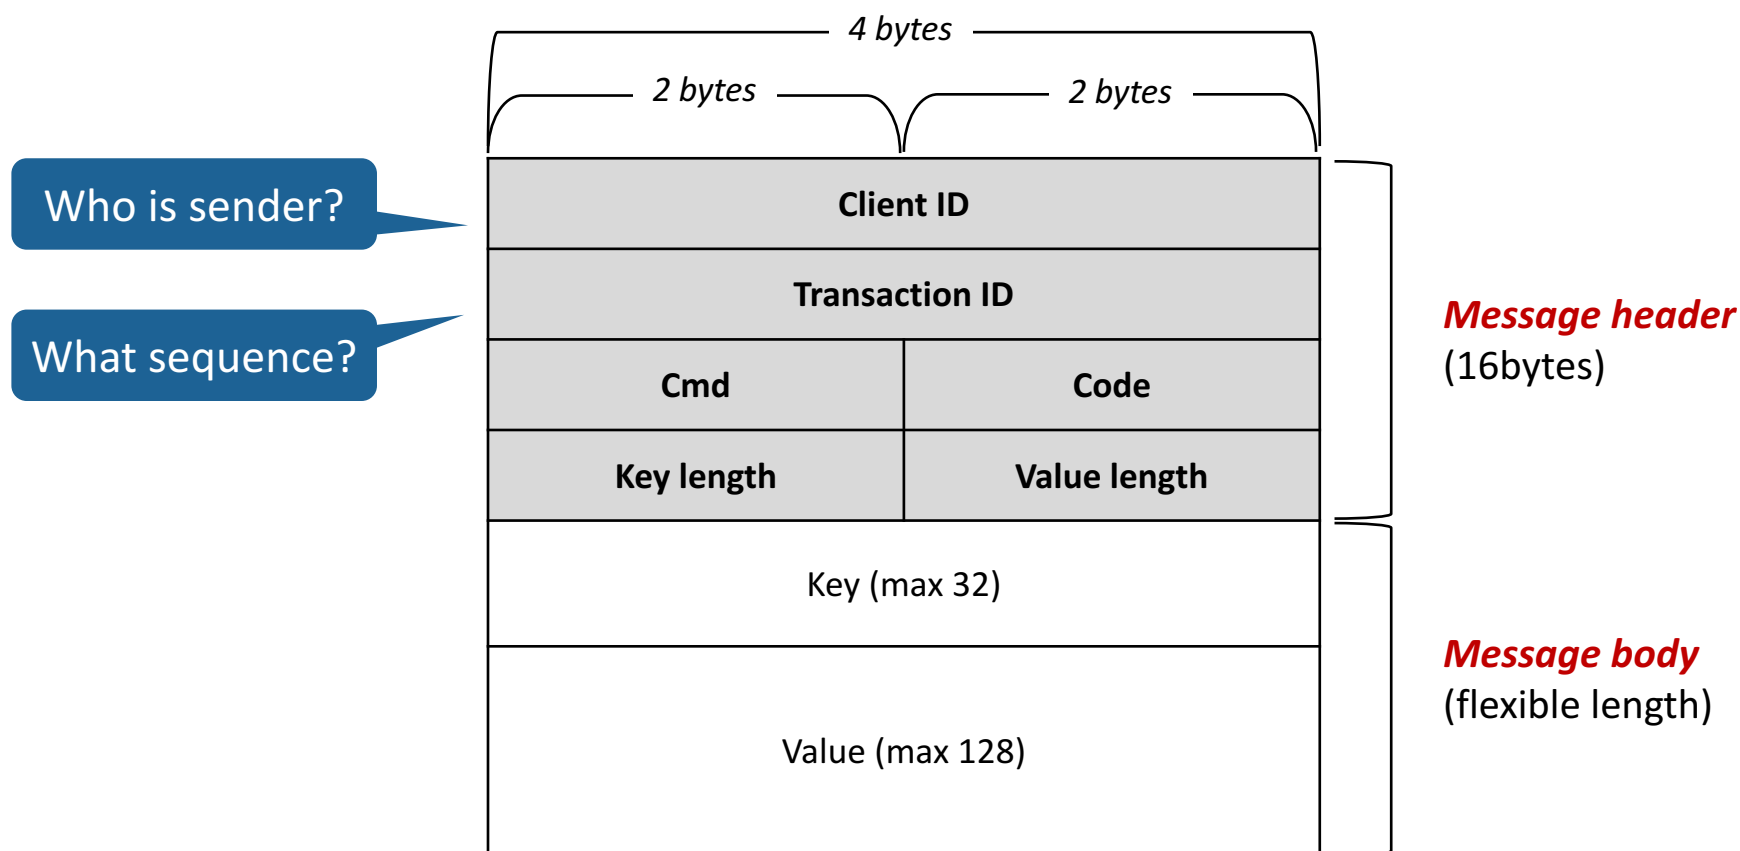
Overall architecture

- All components = processes



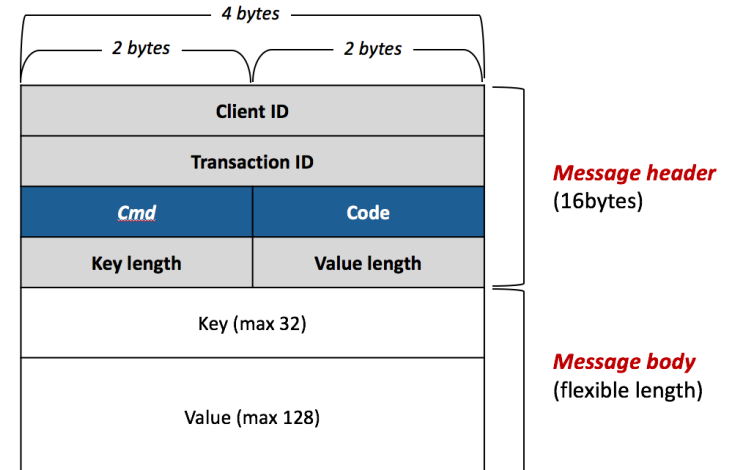
Protocol

- Data format



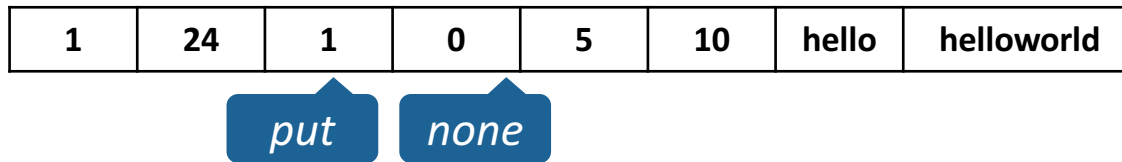
Protocol(Cont'd)

- Cmd
 - Identify **commands**
 - **put**: 1, put-ack: 2, **get**: 3, get-ack: 4, **del**: 5, del-ack: 6
 - Code
 - Represent **status** for put-ack, get-ack, and del-ack
 - **none***: 0, success: 1, not exist: 2, already exist: 3
- * Only used for non-ack messages

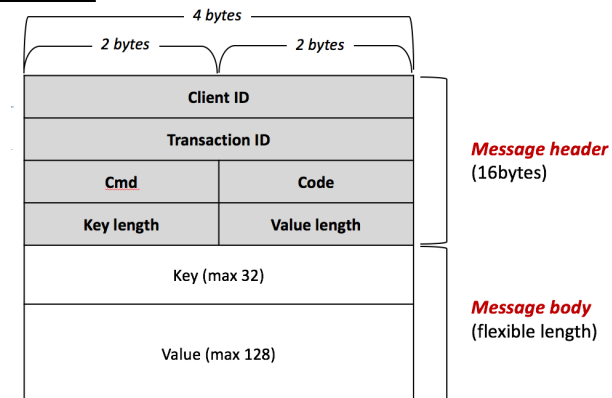
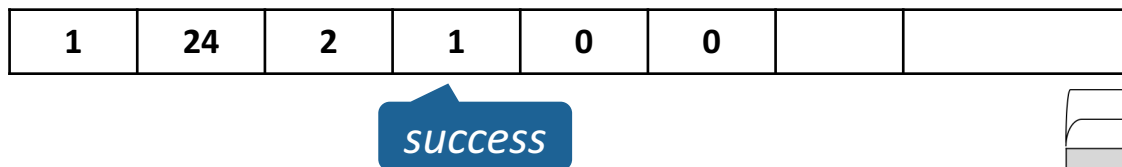


Examples

- Client #1 requests **“put”** a {hello, helloworld} pair into the system (**sequence #24**)

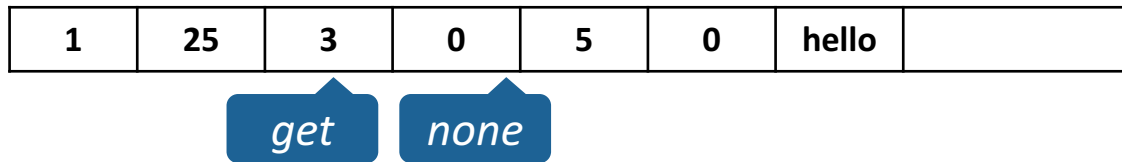


- The **“put”** message requested by Client #1 was successfully stored on the system (**sequence #24**)

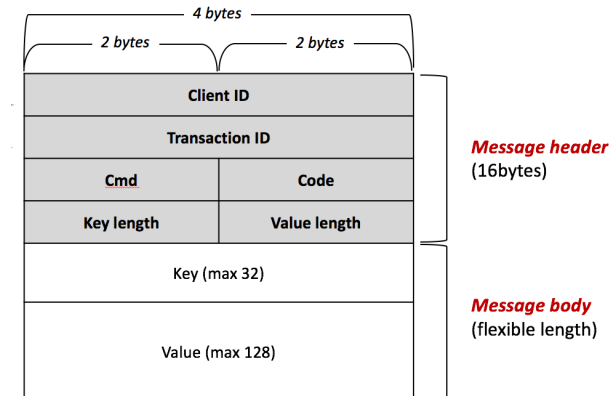
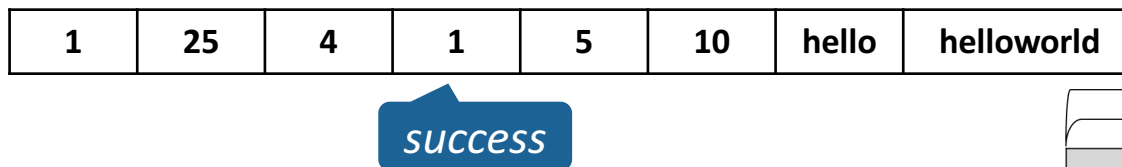


Examples(Cont'd)

- Client #2 requests **“get”** a {hello} key to the system (**sequence #25**)

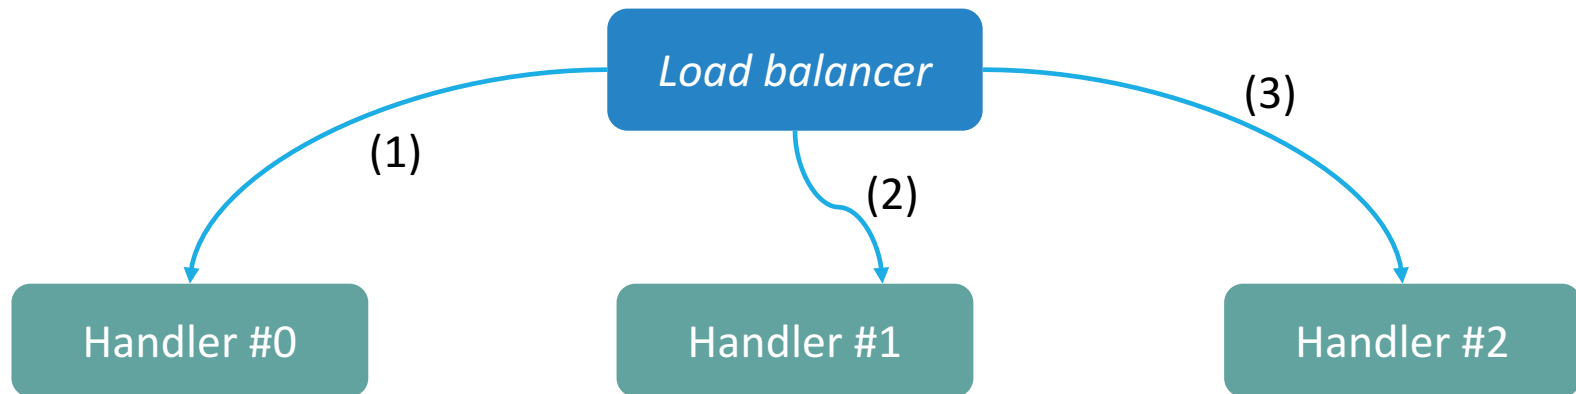


- The **“get”** message requested by Client #2 was successfully found on the system(**sequence #25**)



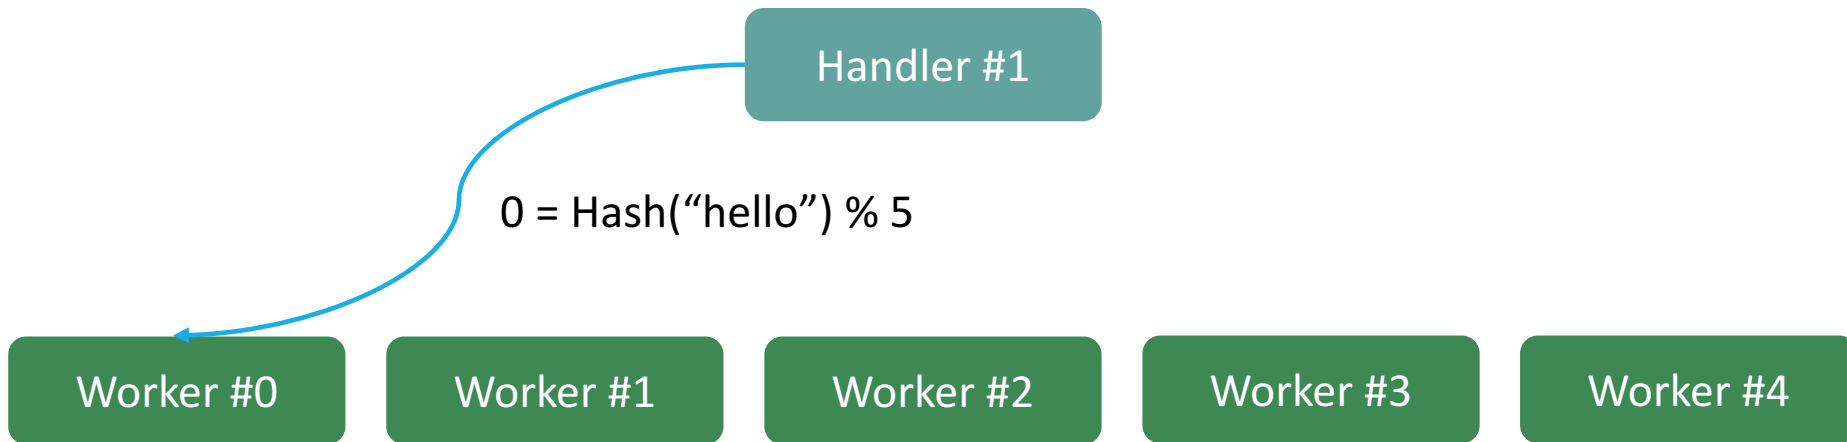
Operations – Load balancer

- Request delivery to *a specific handler* according to a LB strategy
➔ *Round-robin*
- Response delivery to the corresponding client



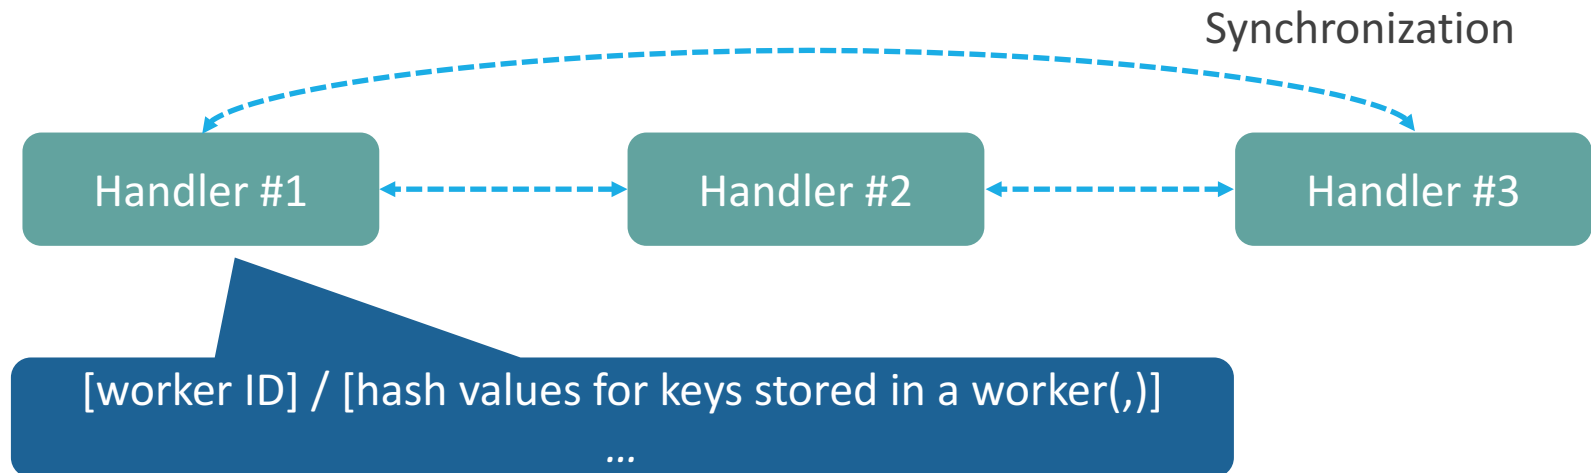
Operations - Handler

- Request delivery to a specific worker and response delivery to the load balancer
 - **Target worker = Hash(key) % # of workers**
 - You should use **Jenkins hash function(one-at-a-time)**
 - See https://en.wikipedia.org/wiki/Jenkins_hash_function



Operations – Handler(Cont'd)

- Management of **key-value distributions**
 - **Synchronization** of key-value distributions among handlers
 - All handlers should eventually know all key-value distributions
- * ignore the strong consistency and failure issues*



Operations - Worker

- **Put** a pair of a key and a value into its data table
- **Get/Del** a value of a key from its data table

** ignore the failure issue*

Worker #1

Worker #2

Worker #3

Worker #4

Worker #5

[hash value of a key] / [key] / [value]

...

[hash value of a key] / [key] / [value]

...

[hash value of a key] / [key] / [value]

...

Requirement #1 - Message logging

- All received messages *should be written* in log files
- Format: *[component_name].log* (e.g., handler_2.log, worker_1.log)
- Contents:
 - *[time] [client id] [transaction id] [cmd] [code] [key length] [value length] [key] [value]*
 - *[time] [client id] [transaction id] [cmd] [code]*

Requirement #2 - CLI

- Client's CLI(Example)

```
Client> connect [LB IP]
```

```
Client> set 1
```

```
Client #1>
```

```
Client #1> put [key] [value]
```

```
Success
```

```
or
```

```
Fail (reason: already exist)
```

```
Client #1> get [key]
```

```
[value]
```

```
or
```

```
Fail (reason: not exist)
```

```
Client #1> del [key]
```

```
Success
```

```
or
```

```
Fail (reason: not exist)
```

Requirement #2 – CLI(Cont'd)

- Load-balancer's CLI

```
LB> list  
[handler ID] / [handler IP address / port] / [# of requests]
```

- Handler's CLI

```
HD# [handler ID]> list  
[worker ID] / [hash values for keys stored in a worker(,)]  
HD# [handler ID]> show [worker ID]  
[hash values for keys stored in the worker]
```

- Worker's CLI

```
WK# [worker ID]> list  
[hash value of a key] / [key] / [value]  
WK# [worker ID]> show [key]  
[hash value of the key] / [key] / [value]
```

Notice

- Languages
 - C / Java
 - ***Please DO write comments in your code (be concise)***
- Tests
 - Your system MUST be implemented as ***separate modules***(Clients, Load Balancer, Handlers, Workers)
 - TAs will conduct ***module tests*** to evaluate your system
 - TAs will validate your system ***with other students' clients***
- Environments
 - Ubuntu 14.04.5 LTS (Trusty Tahr) 64-bit

Notice(Cont'd)

- We **highly** recommend that you use **vagrant**
 - Install VirtualBox: <https://www.virtualbox.org/>
 - Install Vagrant: <https://www.vagrantup.com/downloads.html>
 - Type the following commands
 - \$ mkdir PATH_TO_BOX; cd PATH_TO_BOX
 - \$ vagrant init ubuntu/trusty64
 - \$ vagrant up
 - \$ vagrant ssh

} You just need to use this after installation

** Borrowed by Sang Kil Cha's slide*

Grade

- Total **200pts**
 - **40pts** Documentation (Including README)
 - **20pts** Request and response in a client
 - **20pts** Load-balance requests to handlers in a load balancer
 - **20pts** Maintain key-value distributions in a handler
 - **20pts** Synchronize the distributions among handlers
 - **20pts** Manage a key-value table in a worker

Grade(Cont'd)

- 60pts *handle burst requests*
 - 10,000 (10pts) ~ 100,000 (60pts) requests per second
 - 100,000 *random unique* key-value requests
 - **Put** 50,000 pairs
 - **Get** 25,000 pairs (1/2 pairs exist), and **Del** 25,000 pairs (1/2 pairs exist)

